

An Introduction to Graph Databases

Simon Larsén - slarse@kth.se

DD2482 DevOps and Automated Testing

2019-04-27

Contents

1	Introduction	1
2	Background	1
2.1	Resource Description Framework (RDF) databases	1
2.2	Property graph databases	2
2.3	A note on performance	3
3	Usage examples: SQL versus PG	3
3.1	SQL database definition	3
3.2	Property graph database definition	4
3.3	Querying the databases	5
3.3.1	Query #1: Find all actors and the roles they have played	5
3.3.2	Query #2: Find all self-directed actors	6
3.3.3	Query #3: Rumors about Ben	6
4	Discussion	8
5	References	10
6	Appendix	11
6.1	Appendix A	11
6.2	Appendix B	12
6.3	Appendix C	13

1 Introduction

For decades, the word database has been more or less synonymous with relational database, which in turn is more or less synonymous with SQL database. There is no denying that SQL database systems are both performant and reliable, and have served many purposes very well. However, using a single technology to tackle all storage-related problems fundamentally contradicts a simple intuition about software engineering: there is no one size fits all solution. This article presents an introduction to graph databases, which is a different take on data storage that does not build on the relational model. Readers are expected to at least have a cursory understanding of SQL and the relational model, as these will be used for comparison without being explained in detail. The meat of the article is a theoretical presentation of two common graph database models, as well as a by-example usage comparison of the Neo4j graph database system with the PostgreSQL relational database system. After having read this article, a reader with some previous experience of using SQL databases should walk away with some rudimentary knowledge of when it may be appropriate to use a graph database instead, and how to get started doing so.

2 Background

Most real world applications need persistent data storage, and for a long time, the relational database model has been the go-to for most such data stores [1]. The 1970s saw the inception of the *Standard Query Language* (SQL), as well as seminal work on performance optimizations that paved the way for relational databases in the world outside of academia [2]. Adoption rose sharply in the 1980s and SQL-based databases has since been the de-facto standard. In the past few years, there has been an increase in interest for non-relational database systems, collectively known as NoSQL [1]. It is important to note that NoSQL is not *one* technology, but rather denotes a collection of storage technologies that are not based on the relational database model. Graph databases form a subset of the NoSQL technologies [3]. The concept of graph databases has been around since the 1980s, but was more or less forgotten about in the mid 1990s due to the emergence of other new technologies, such as XML. It is not until recently that graph databases have seen a resurgence in popularity. Much due to this fact, the field is still in flux, and there are many different types of graph database models around. Angles et al. provides a fairly comprehensive but slightly dated overview of several different models and their origins [4]. As there is such a large amount of different graph database models, this article will not cover all of them. Rather, I will introduce the *property graph* database, and show by practical example how using such a database compares to the use of a traditional SQL database. I will also present the concepts of a different kind of graph database based on the RDF standards, as such databases frequently appeared during background research of the subject [5], [6]. The point of including some theory on RDF graph databases is mostly to show that just as NoSQL is not one technology, graph databases is not one technology. The focus of the rest of the article is on property graphs. Note also that in this article, the term graph database is used interchangeably with property graph database, and any other type of graph database is referred to by a different name. The rest of the article is structured as follows. Section 2.1 presents the concepts behind RDF databases, while Sec. 2.2 does the same for property graph databases. Section 2.3 makes a brief detour into the complicated area of performance measurements. Section 3 compares creation and subsequent querying of a movie database using the relational PostgreSQL system, and the property graph Neo4j system. Finally, Section 4 presents a discussion of the potential pros and cons of using a graph database over a relational database.

2.1 Resource Description Framework (RDF) databases

The *Resource Description Framework* (RDF) is a set of specifications for expressing metadata for websites, and is a key part of the WWW Consortium's (W3C) effort to standardize a Semantic Web that a machine can not only navigate, but also understand [7], [8]. Lately, it has also seen use as general purpose data storage in several products, including AllegroGraph¹ and GraphDB². An RDF graph is a set of triples of *subject*, *predicate* and *object*. The subject is that which we want to say something about, the predicate is what kind of statement we are making, and the object the value of that statement. For

¹<https://franz.com/agraph/allegrograph/>

²<http://graphdb.ontotext.com/>

example, the triples in Table 1 is an RDF-like encoding of the data “The Town is a movie”, “Ben is a person”, “Ben acted in The Town”, and “The Town was directed by Ben”.

Table 1: RDF triples example.

Subject	Predicate	Object
The Town	type	Movie
Ben	type	Person
Ben	actedIn	The Town
The Town	directedBy	Ben

There is an interesting observation to be made regarding the triples presented in Table 1: it is not immediately apparent that they constitute a graph. In fact, it is possible to have for example the object of one triple be the predicate of another, so the triples do not by themselves represent a graph in the traditional sense [9]. In the case of Table 1 however, viewing the set union of subjects and objects as nodes, and the set of predicates as labeled edges, it is clear that the triples of Table 1 induce the graph in Fig. 1.




Figure 1: Visualisation of Table 1

Another interesting observation is that Table 1 looks similar to a traditional relational database table, and indeed, there are ways to implement RDF stores on top of a relational database system [10].

2.2 Property graph databases

Unlike RDF graphs, property graphs (PG) are not formally defined [5]. The term also seems to be fairly new; the earliest mention that I can find is from 2010 [11], while another contemporary author does not use the term to describe the same concept [12]. There is however a general agreement about what a property graph is: a directed graph in which nodes and edges are distinct, labeled entities that can contain properties in the form of key-value pairs. Note that the edges of a PG are often referred to as *relationships*, but I will refer to them as edges. Slightly adding to the confusion is that the term property graph is often used interchangeably with *labeled property graph* (LPG) [5], [13], [14]. The L in LPG simply emphasizes the fact that nodes and edges are labeled, which is sensible considering the important role that labels play in a PG model. Labels are used to categorize nodes and edges, and are roughly equivalent to type assignments [12]. For example, a node representing a person could have the label “Person”, which makes it easy to query for Person node. Figure 2 shows a property graph example of the same data that was presented in Fig. 1.

Another important difference between RDF graphs and property graphs is that the latter has no standardized query language, although it should be noted that efforts are underway to standardize




Figure 2: Visualisation of a property graph. Labels are in bold and properties are written **key: value**

a *Graph Query Language* (GQL) [6], [15]–[17]. Examples of current PG databases are Neo4j³, which uses the Cypher query language⁴ (open-sourced as openCypher⁵), and JanusGraph⁶, which uses the Gremlin query language⁷.

2.3 A note on performance

A topic which this article has conveniently avoided thus far is that of performance. The reason is simple: one could write an entire article series only on that, as performance comparisons are very difficult to perform between systems that are designed in fundamentally different ways. There is however research on the topic, both comparing graph databases to other graph databases [18]–[20], and comparing relational databases to graph databases [1]. The only result that I find relevant to mention for this article is that a graph database system such as Neo4j can perform on par with, or even outperform, a relational SQL database system in some cases, but also falls well behind in others [1]. This means that a graph database can be a viable alternative to a relational database even for performance critical applications, but it may not always be the right choice.

3 Usage examples: SQL versus PG

In this section, I present a simple movie dataset modelled with a property graph database (Neo4j) and a standard relational database (PostgreSQL). The dataset consists of the following:

- People
- Movies
- Which movies any given person has acted in
 - And which role he or she performed as
- Which movies any given person has directed

Modeling this data represents a few common tasks database systems have to fulfill, namely representing concrete entities (here, movies and people) and their attributes (names and titles), as well as relationships between such entities (“person acted in movie” and “movie was directed by person”). The rest of this section is structured as follows. Section 3.1 presents how to define and populate a database with SQL, and Sec. 3.2 does the same for a property graph database. Section 3.3 then presents a series of queries on the databases.

3.1 SQL database definition

This is the part of the article that assumes some prior knowledge of relational databases, as concepts such as tables and foreign keys will not be explained. To model the data mentioned in Sec. 3, a

³<https://neo4j.com>

⁴<https://neo4j.com/developer/cypher/>

⁵<http://www.opencypher.org/>

⁶<https://janusgraph.org/>

⁷<https://docs.janusgraph.org/latest/gremlin.html>

typical SQL database will need four tables: two tables to represent the base entities **Person** and **Movie**, as well as two association tables to model the **ActedIn** and **DirectedBy** relationships between these. The reason that the two association tables are required is that both **ActedIn** and **DirectedBy** are many-to-many relationships, which cannot be expressed in a relational database with only the concrete entity tables. For example, an actor typically has acted in more than one movie, and most movies are acted in by more than one actor. Listing 1 shows an example of how the **Person** table could be defined, with the **Movie** table being almost identical. The **ActedIn** table could then be defined as shown in Listing 2, with the **DirectedBy** table being almost identical to that.

Listing 1 DDL for the **Person** table

```
CREATE TABLE Person (
  id INT UNSIGNED PRIMARY KEY AUTO_INCREMENT,
  name VARCHAR(128) NOT NULL
);
```

Listing 2 DDL for the **ActedIn** association table

```
CREATE TABLE ActedIn (
  id SERIAL PRIMARY KEY,
  person_id INT NOT NULL REFERENCES Person(id),
  movie_id INT NOT NULL REFERENCES Movie(id),
  played_role VARCHAR(128) NOT NULL
);
```

The whole database schema is presented schematically in Fig. 3. Note the multiplicities on the relations between the tables. For example, each **ActedIn** row is associated with precisely one **Person** row, while each **Person** row is associated with zero (because not every person is an actor) or more **ActedIn** rows.

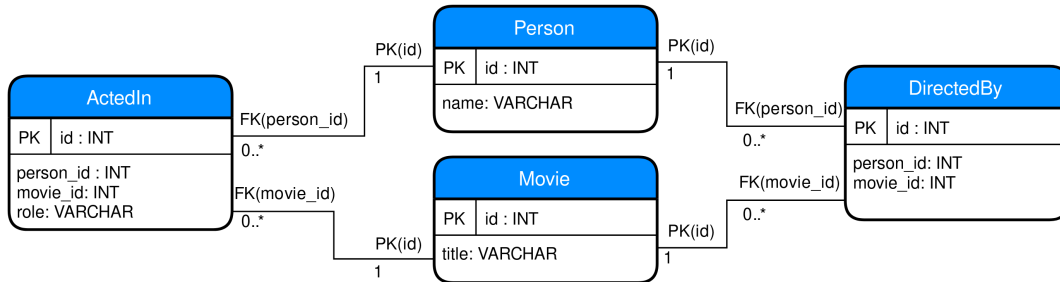


Figure 3: SQL schema for the movie database

With the schema defined, the database can be filled with data using queries like those shown in Listing 3. Note how the insertions into the association tables have nested **SELECT** queries to find the correct primary keys of the related tables. This is a bit cumbersome, but I know of no other standard way to achieve it. The full definition of the database, including data insertions, can be found in Sec. 6.1. It has been tested to work with PostgreSQL 9.6.

3.2 Property graph database definition

For the property graph database, I use Neo4j and its query language Cypher. The reason is mostly that it is easy to get started with Neo4j, and Cypher is easy to briefly explain. Do however keep in mind that there is no one query language for property graphs, as described in Sec. 2.2, so this section is not representative of property graphs as a whole. I do however think that it illustrates the idea of graph-based queries well. As Neo4j is a schemaless database system, there is no need to first define the database, as was the case for the SQL database in Sec. 3.1. It is simply a matter of entering values into the database. Cypher is concise, so entering the same data about Ben Affleck that was entered in Sec. 3.1 is a matter of the four lines shown in Listing 4.

Listing 3 Sample of value insertions into the SQL database

```
INSERT INTO Movie(title) VALUES
('The Town');
INSERT INTO Person(name) VALUES
('Ben Affleck');
INSERT INTO ActedIn(person_id, movie_id, played_role) VALUES
((SELECT id FROM Person WHERE name='Ben Affleck'),
 (SELECT id FROM Movie WHERE title='The Town'),
 'Doug MacRay');
INSERT INTO DirectedBy(person_id, movie_id) VALUES
((SELECT id FROM Person WHERE name='Ben Affleck'),
 (SELECT id FROM Movie WHERE title='The Town'));
```

Listing 4 Sample CREATE statement for the Neo4j database

```
CREATE (thetown:Movie {title: "The Town"})
CREATE (ben:Person {name: "Ben Affleck"})
CREATE (ben)-[:ACTED_IN {role: "Doug MacRay"}]-(thetown)
CREATE (thetown)-[:DIRECTED_BY]-(ben)
```

Listing 4 is technically one statement, and the three **CREATE**s following the first could be replaced with commas to make it even more concise (see the full database definition in Sec. 6.2 for an example). The first two lines create **Person** and **Movie** nodes, while the last two create **ACTED_IN** and **DIRECTED_BY** edges (or relations). The syntax for the **CREATE** statement is rather straightforward. Creating a node is as simple as **CREATE (<NODE_DEFINITION>)**. Looking specifically at the first line, **thetown:Movie** denotes that the label on this node is **Movie**, and that the node should be assigned to a variable called **thetown**. The variable can be used throughout the query to reference the node. Finally, everything within curly braces are simply properties on the form **key: value**. Neo4j has an excellent visualisation tool built-in to its web interface, and the visualisation of this particular database can be found in Fig. 4.

The definition of a relationship is similarly straightforward, and generally looks like **CREATE (<NODE>)-[<EDGE_DEFINITION>]->[<NODE>]**. **<NODE>** can be for example a variable or a node definition, and the **<EDGE_DEFINITION>** is written precisely the same as **<NODE_DEFINITION>**, including property declarations. Looking more closely at the last **CREATE** statement, it is plain to see how the variables **thetown** and **ben** are reused to denote the nodes which the edge connects. Note the arrow at one end of the edge defining its direction (**->**). All edges in a Neo4j database are directed, but Sec. 3.3.3 will show that queries can treat them as if they were undirected. Also note how no variable name is declared for the edge. There is no need for it, as it is not reused. The full database definition can be found in Sec. 6.2.

3.3 Querying the databases

This section presents a series of increasingly complex queries in SQL and Cypher. Keep in mind that these queries play to the strengths of a graph database. The idea is to show that use cases where graph databases may be preferable to relational databases *exist*, and not that they are preferable as a general rule. The expected results of all queries can be found in Sec. 6.3.

3.3.1 Query #1: Find all actors and the roles they have played

For the first query, we want to list all actors and the roles they have played, as well as in which movies. This is a straightforward three-way join with SQL, as shown in Listing 5.

With Cypher, it gets slightly simpler, using the **MATCH** statement shown in Listing 6. A **MATCH** statement looks much like a **CREATE** statement, but instead of declaring a structure to create, it declares a structure to search for. Here, that structure is any **Person** node connected to a **Movie** node via an **ACTED_IN** edge. Note that the edge is directed. Note how the nodes are assigned to variables **person** and **movie**, and the edge to **acted**, and that these variables are then used to select the specific results required

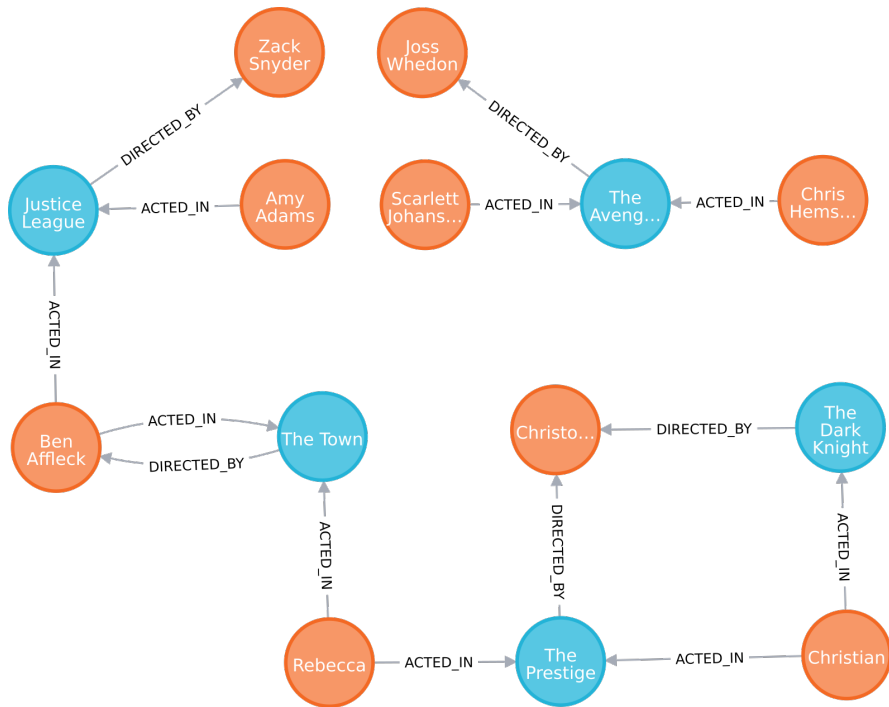


Figure 4: Graph visualisation of the example database. Orange nodes represent people, and blue nodes represent movies.

Listing 5 SQL query #1

```

SELECT Person.name, ActedIn.played_role, Movie.title
FROM Person, Movie, ActedIn
WHERE
  Person.id = ActedIn.person_id AND
  Movie.id = ActedIn.movie_id;

```

in the **RETURN** statement. The results of the SQL and Cypher queries should be identical, apart from potentially different ordering of the results

Listing 6 Cypher query #1

```

MATCH (actor:Person)-[acted:ACTED_IN]->(movie:Movie)
RETURN actor.name, acted.played_role, movie.title

```

3.3.2 Query #2: Find all self-directed actors

This query is meant to find all actors that have acted in a movie that they have also directed. For SQL, this results in a slight extension of the three-way join in Sec. 3.3.1, making it the four-way join shown in Listing 7.

The Cypher query also needs an extension, but it is done in a different way by simply adding another structure to match against. The query can be found in Listing 8, which clearly shows how convenient Cypher's variables can be. Because the same variables appear in both structures, the comma can be viewed as a logical AND: only **Person** and **Movie** nodes connected by both an **ACTED_IN** edge and a **DIRECTED_BY** edge are matched. Again, both SQL and Cypher queries will yield identical results.

3.3.3 Query #3: Rumors about Ben

This is the final query, which is meant to clearly demonstrate the benefit of reasoning about data as a graph. The idea is to find any actor that may have heard a rumor about what it is like acting together

Listing 7 SQL query #2

```
SELECT Person.name, ActedIn.played_role, Movie.title
FROM Person, Movie, ActedIn, DirectedBy
WHERE
    Person.id = ActedIn.person_id AND
    Person.id = DirectedBy.person_id AND
    Movie.id = ActedIn.movie_id AND
    Movie.id = DirectedBy.movie_id;
```

Listing 8 Cypher query #2

```
MATCH (actor:Person)-[acted:ACTED_IN]->(movie:Movie),
      (movie)-[:DIRECTED_BY]->(actor)
RETURN actor.name, acted.played_role, movie.title
```

with Ben Affleck. This includes any actor A who has acted in the same movie as Ben, or any actor B who has acted with actor A, or any actor C who has acted with actor B, and so on. The astute reader may have noticed that this is a slightly more complicated version of traversing a social graph to find a person’s friends, friends of friends, and so on⁸. The problem can be somewhat simplified by considering the DIRECTED_BY and ACTED_IN edges visualised in Fig. 4 as undirected edges. Then, it is simply a matter of finding every actor that can be reached from Ben’s node by traversing ACTED_IN edges. While it does not make much sense for a movie to have acted in an actor, it makes the problem easier to reason about. To solve such a graph traversal problem in SQL, we need to issue a so called *hierarchical query*⁹, shown in Listing 9.

Listing 9 SQL query #3

```
WITH RECURSIVE acted_in(person_id, movie_id) AS (
    /* Initial start query */
    SELECT person_id, movie_id
    FROM ActedIn, Person
    WHERE Person.id = ActedIn.person_id AND
           Person.name = 'Ben Affleck'
UNION
    /* The recursive query */
    SELECT ActedIn.person_id, ActedIn.movie_id
    FROM acted_in, ActedIn
    WHERE acted_in.movie_id = ActedIn.movie_id OR
           acted_in.person_id = ActedIn.person_id
)
/* selecting from the result */
SELECT DISTINCT Person.name
FROM acted_in, Person
WHERE acted_in.person_id = Person.id
AND Person.name != 'Ben Affleck'
```

I will briefly explain what is actually happening in the query. On the first line, the `acted_in` pseudo-entity is defined. Then, an initial “start” query is issued to find all of Ben’s (`person_id`, `movie_id`) tuples. What happens next is fairly unintuitive. The results of the “recursive” query is unioned with the initial query, the result of which is then unioned with the recursive query again, and again, until the result set is no longer expanding. It is essentially a breadth first search over the `ActedIn` relations, where the movies are discarded in the final result. The Cypher equivalent of the SQL query shown in Listing 10 is a good example of why representing data as graphs can be advantageous.

⁸This is closely related to the transitive closure of a binary relation R on some set S. Wikipedia has a nice page on the subject: https://en.wikipedia.org/wiki/Transitive_closure

⁹Again, Wikipedia has a nice page on the subject: https://en.wikipedia.org/wiki/Hierarchical_and_recursive_queries_in_SQL

Listing 10 Cypher query #3

```
MATCH (ben:Person {name: "Ben Affleck"}),  
      (ben)-[:ACTED_IN*]-(movie:Movie),  
      (actor:Person)-[:ACTED_IN]->(movie)  
RETURN actor.name
```

Fetching the Ben Affleck node and storing it in the `ben` variable is not strictly necessary, and could be done inline on the second line, but I found this more readable. Listing 10 can be broken down as follows.

1. `MATCH (ben:Person {name: "Ben affleck"})`
 - Find the Ben Affleck node and store it in `ben`
2. `(ben)-[:ACTED_IN*]-(movie:Movie)`
 - Find all `Movie` nodes reachable by traversing any amount of `ACTED_IN` edges, starting from `ben`
 - Note the `*` for “any amount”
 - Note the lack of an arrowhead on the edge, which instructs Neo4j to consider the `ACTED_IN` edges as undirected
3. `(actor:Person)-[:ACTED_IN]->(movie)`
 - Find every actor who acted in any of the found `movie` nodes
4. `RETURN actor.name`
 - Return the name of every actor that matched the constraints

The results of SQL and Cypher queries should yet again be identical.

4 Discussion

Property graph databases show some clear advantages over traditional relational databases. The queries in Cypher are shorter throughout, yet remain readable. The final query exemplified how a non-trivial but still realistic search query required use of advanced SQL constructs, with quite a lot of boilerplate, while the Cypher equivalent only made use of some of the most rudimentary features of the language and remained concise. While shorter is not always better, shorter *and* more readable surely has to count as better from a maintainability point of view. Another stark contrast between SQL and Cypher is the former’s reliance on a pre-defined schema, and the latter’s lack of such. A schemaless approach is a great boon to quick iteration development practices, as many minor alterations to the abstract schema of a schemaless database may not require any production database maintenance. On the other hand, *any* alteration in a SQL schema requires a database migration in production environments, which from personal experience can be a hassle even for relatively simple databases. This schemaless nature does however have the drawback that the data format is in not rigidly defined, which could lead to unexpected runtime behavior because some expected property was simply missing from a node or edge. It could be argued that the same can happen in a relational database by putting null values all over the place, so as with many other things, the advantages and disadvantages of either approach is not clear-cut.

There are some major caveats to these comparisons that should not go without mention. First, I specifically modelled the database and queries to play to the strengths of a graph representation. In other words, these are by no means general conclusions about graph databases being superior to relational databases, but rather that they can be given the right circumstances. Furthermore, Cypher is not *the* property graph database language, so results are not even generalizable over PG query languages as a whole. Time will tell if GQL ends up becoming a widely adopted standard, but as it stands, there is nothing resembling a standard for property graph databases. Adopting a database such as Neo4j may therefore lead to more lock-in than risk averse managers may be willing to accept. My closing thoughts on the matter are however that property graph databases are not here to replace relational databases, but definitely show promise as being superior in some circumstances. This is especially true for an environment where the data is easily modelled as a graph, where graph-like queries are common, or where rapid iteration on the related product may cause the model to be extended or changed over a relatively short amount of time. The fact that a relatively new database system such as Neo4j in some

cases can rival the performance of relational database systems that have been developed and optimized for decades is also encouraging. I think that all of this makes it quite clear that relational databases as the one size fits all solution to persistent data storage is a thing of the past, and that there are now other solutions that sometimes provide a better fit.

5 References

- [1] C. Vicknair, M. Macias, Z. Zhao, X. Nan, Y. Chen, and D. Wilkins, “A comparison of a graph database and a relational database: A data provenance perspective,” in *Proceedings of the 48th annual southeast regional conference*, 2010, p. 42.
- [2] P. G. Selinger, M. M. Astrahan, D. D. Chamberlin, R. A. Lorie, and T. G. Price, “Access path selection in a relational database management system,” in *Proceedings of the 1979 acm sigmod international conference on management of data*, 1979, pp. 23–34.
- [3] M. Buerli and C. Obispo, “The current state of graph databases,” *Department of Computer Science, Cal Poly San Luis Obispo*, *mbuerli@calpoly.edu*, vol. 32, no. 3, pp. 67–83, 2012.
- [4] R. Angles and C. Gutierrez, “Survey of graph database models,” *ACM Computing Surveys (CSUR)*, vol. 40, no. 1, p. 1, 2008.
- [5] O. Hartig, “Reconciliation of rdf* and property graphs,” *arXiv preprint arXiv:1409.3288*, 2014.
- [6] R. Angles, M. Arenas, P. Barceló, P. Boncz, G. Fletcher, C. Gutierrez, T. Lindaaker, M. Paradies, S. Plantikow, J. Sequeda, and others, “G-core: A core for future graph query languages,” in *Proceedings of the 2018 international conference on management of data*, 2018, pp. 1421–1432.
- [7] R. Cyganiak, D. Wood, and M. Lanthaler, “RDF 1.1 concepts and abstract syntax,” 2014. [Online]. Available: <https://www.w3.org/TR/rdf11-concepts/#section-rdf-graph>. [Accessed: 04-Apr-2019].
- [8] “Semantic web.” [Online]. Available: <https://www.w3.org/standards/semanticweb/>. [Accessed: 04-Apr-2019].
- [9] J. Hayes and C. Gutierrez, “Bipartite graphs as intermediate model for rdf,” in *International semantic web conference*, 2004, pp. 47–61.
- [10] M. A. Bornea, J. Dolby, A. Kementsietsidis, K. Srinivas, P. Dantressangle, O. Udrea, and B. Bhattacharjee, “Building an efficient rdf store over a relational database,” in *Proceedings of the 2013 acm sigmod international conference on management of data*, 2013, pp. 121–132.
- [11] M. A. Rodriguez and P. Neubauer, “Constructions from dots and lines,” *Bulletin of the American Society for Information Science and Technology*, vol. 36, no. 6, pp. 35–41, 2010.
- [12] S. Srinivasa, “Data, storage and index models for graph databases,” in *Graph data management: Techniques and applications*, IGI Global, 2012, pp. 47–70.
- [13] “RDF triple stores vs. Labeled property graphs: What’s the difference.” [Online]. Available: <https://neo4j.com/blog/rdf-triple-store-vs-labeled-property-graph-difference/>. [Accessed: 02-Apr-2019].
- [14] “Graph database.” [Online]. Available: https://en.wikipedia.org/wiki/Graph_database. [Accessed: 02-Apr-2019].
- [15] “GQL standrd,” 2018. [Online]. Available: <https://www.gqlstandards.org/>. [Accessed: 04-Apr-2019].
- [16] “The gql manifesto,” 2018. [Online]. Available: <https://gql.today/>. [Accessed: 04-Apr-2019].
- [17] D. Ragget, “W3C workshop on web standardization for graph data,” 2019. [Online]. Available: <https://www.w3.org/Data/events/data-ws-2019/>. [Accessed: 04-Apr-2019].
- [18] R. C. McColl, D. Ediger, J. Poovey, D. Campbell, and D. A. Bader, “A performance evaluation of open source graph databases,” in *Proceedings of the first workshop on parallel programming for analytics applications*, 2014, pp. 11–18.
- [19] S. Jouili and V. Vansteenbergh, “An empirical comparison of graph databases,” in *2013 international conference on social computing*, 2013, pp. 708–715.
- [20] M. Ciglan, A. Averbuch, and L. Hluchy, “Benchmarking traversal operations over graph databases,” in *2012 ieee 28th international conference on data engineering workshops*, 2012, pp. 186–189.

6 Appendix

6.1 Appendix A

```
CREATE TABLE Person (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(128) NOT NULL  
);  
  
CREATE TABLE Movie (  
    id SERIAL PRIMARY KEY,  
    title VARCHAR(128) NOT NULL  
);  
  
CREATE TABLE DirectedBy (  
    id SERIAL PRIMARY KEY,  
    person_id INT NOT NULL REFERENCES Person(id),  
    movie_id INT NOT NULL REFERENCES Movie(id)  
);  
  
CREATE TABLE ActedIn (  
    id SERIAL PRIMARY KEY,  
    person_id INT NOT NULL REFERENCES Person(id),  
    movie_id INT NOT NULL REFERENCES Movie(id),  
    played_role VARCHAR(128) NOT NULL  
);  
  
INSERT INTO Movie(title) VALUES  
    ('The Avengers'),  
    ('The Town'),  
    ('Justice League'),  
    ('The Prestige'),  
    ('The Dark Knight');  
  
INSERT INTO Person(name) VALUES  
    ('Amy Adams'),  
    ('Ben Affleck'),  
    ('Chris Hemsworth'),  
    ('Scarlett Johansson'),  
    ('Rebecca Hall'),  
    ('Christian Bale'),  
    ('Christopher Nolan'),  
    ('Zack Snyder'),  
    ('Joss Whedon');  
  
INSERT INTO ActedIn(person_id, movie_id, played_role) VALUES  
    ((SELECT id FROM Person WHERE name='Amy Adams'),  
     (SELECT id FROM Movie WHERE title='Justice League'),  
     'Lois Lane'),  
    ((SELECT id FROM Person WHERE name='Ben Affleck'),  
     (SELECT id FROM Movie WHERE title='The Town'),  
     'Doug MacRay'),  
    ((SELECT id FROM Person WHERE name='Ben Affleck'),  
     (SELECT id FROM Movie WHERE title='Justice League'),  
     'Batman'),  
    ((SELECT id FROM Person WHERE name='Chris Hemsworth'),  
     (SELECT id FROM Movie WHERE title='The Avengers'),
```

```

    'Thor'),
  ((SELECT id FROM Person WHERE name='Christian Bale'),
   (SELECT id FROM Movie WHERE title='The Prestige'),
   'Alfred Borden'),
  ((SELECT id FROM Person WHERE name='Christian Bale'),
   (SELECT id FROM Movie WHERE title='The Dark Knight'),
   'Batman'),
  ((SELECT id FROM Person WHERE name='Rebecca Hall'),
   (SELECT id FROM Movie WHERE title='The Prestige'),
   'Sarah'),
  ((SELECT id FROM Person WHERE name='Rebecca Hall'),
   (SELECT id FROM Movie WHERE title='The Town'),
   'Claire Keesey'),
  ((SELECT id FROM Person WHERE name='Scarlett Johansson'),
   (SELECT id FROM Movie WHERE title='The Avengers'),
   'Black Widow');

INSERT INTO DirectedBy(person_id, movie_id) VALUES
  ((SELECT id FROM Person WHERE name='Ben Affleck'),
   (SELECT id FROM Movie WHERE title='The Town')),
  ((SELECT id FROM Person WHERE name='Christopher Nolan'),
   (SELECT id FROM Movie WHERE title='The Prestige')),
  ((SELECT id FROM Person WHERE name='Christopher Nolan'),
   (SELECT id FROM Movie WHERE title='The Dark Knight')),
  ((SELECT id FROM Person WHERE name='Zack Snyder'),
   (SELECT id FROM Movie WHERE title='Justice League')),
  ((SELECT id FROM Person WHERE name='Joss Whedon'),
   (SELECT id FROM Movie WHERE title='The Avengers'));

```

6.2 Appendix B

```

// Note that separating node/edge definitions with ',' is shorthand
// for multiple CREATE
CREATE (theavengers:Movie {title: "The Avengers"}),
      (thetown:Movie {title: "The Town"}),
      (justiceleague:Movie {title: "Justice League"}),
      (theprestige:Movie {title: "The Prestige"}),
      (thedarkknight:Movie {title: "The Dark Knight"}),
      (amy:Person {name: "Amy Adams"}),
      (ben:Person {name: "Ben Affleck"}),
      (chris:Person {name: "Chris Hemsworth"}),
      (scarlett:Person {name: "Scarlett Johansson"}),
      (rebecca:Person {name: "Rebecca Hall"}),
      (christian:Person {name: "Christian Bale"}),
      (christopher:Person {name: "Christopher Nolan"}),
      (zack:Person {name: "Zack Snyder"}),
      (joss:Person {name: "Joss Whedon"}),
      (amy)-[:ACTED_IN {played_role: "Lois Lane"}]->(justiceleague),
      (ben)-[:ACTED_IN {played_role: "Doug MacRay"}]->(thetown),
      (ben)-[:ACTED_IN {played_role: "Batman"}]->(justiceleague),
      (chris)-[:ACTED_IN {played_role: "Thor"}]->(theavengers),
      (christian)-[:ACTED_IN {played_role: "Alfred Borden"}]->(theprestige),
      (christian)-[:ACTED_IN {played_role: "Batman"}]->(thedarkknight),
      (rebecca)-[:ACTED_IN {played_role: "Sarah"}]->(theprestige),
      (rebecca)-[:ACTED_IN {played_role: "Claire Keesey"}]->(thetown),
      (scarlett)-[:ACTED_IN {played_role: "Black Widow"}]->(theavengers),
      (thetown)-[:DIRECTED_BY]->(ben),

```

```

(theprestige)-[:DIRECTED_BY]->(christopher),
(thedarkknight)-[:DIRECTED_BY]->(christopher),
(justiceleague)-[:DIRECTED_BY]->(zack),
(theavengers)-[:DIRECTED_BY]->(joss)

// Typically, one would also create these indices to boost performance
// NOTE: One query per CREATE INDEX statement, or it will not work

// Index query 1
CREATE INDEX ON :Person(name)

// Index query 2
CREATE INDEX ON :Movie(title)

```

6.3 Appendix C

Table 2: Expected results for query #1

name	played_role	title
Amy Adams	Lois Lane	Justice League
Ben Affleck	Doug MacRay	The Town
Ben Affleck	Batman	Justice League
Chris Hemsworth	Thor	The Avengers
Christian Bale	Alfred Borden	The Prestige
Christian Bale	Batman	The Dark Knight
Rebecca Hall	Sarah	The Prestige
Rebecca Hall	Claire Keesey	The Town
Scarlett Johansson	Black Widow	The Avengers

Table 3: Expected results for query #2

name	played_role	title
Ben Affleck	Doug MacRay	The Town

Table 4: Expected results for query #3

name
Rebecca Hall
Christian Bale
Amy Adams