

KTH DevOps Essay
Continuous Integration and the Cloud

Kasper Liu, Gunnar Applelid

April 2020

Contents

1	Introduction	3
2	Continuous Integration	3
2.1	Jenkins	3
2.2	Codship	4
2.3	Travis	4
2.4	Bamboo	4
2.5	On-Premise CI vs Hosted CI	4
3	Microsoft, Amazon and Google and the "Cloud"	5
3.1	AWS	6
3.2	Google Cloud	6
3.3	Microsoft Azure	6
3.4	Comparison	7
4	Conclusion	7

1 Introduction

The Background of Continuous Integration (CI) is that a small software project with one developer could be an "easy" non-complex process. But what if the number of developers increases and the scope of the software project gets larger? It would then be dangerous to test all the different parts from the developers, just before the product should be delivered to the customer. The risk of problems in the final integration is very high and this could lead to delays in delivering. Therefore, the process of Continuous Integration requires testing of the different components of a software project, many times during the development cycle. This could be done by each developer having their own build which they work on. When they are done with their work, they commit their code to a version control repository. Then an Integration build is conducted together with tests that must pass for the build to be accepted. This process could be done automatically and therefore increase both effectiveness and correctness[1]. It is important to understand that the process itself is what is called Continuous Integration.

2 Continuous Integration

Today there are many tools to aid with continuous integration for all scales of projects. Universities and companies push their employees to develop with CI in mind. There are mainly two types of CI, On-premise CI and hosted CI. Jenkins is a good example of a typical on-premise CI and Codeship offer a service for hosted CI. With the advent of cloud and cloud computing a lot of CI tools have been adopted to cloud environment as well.

2.1 Jenkins

"Jenkins is a free and open source automation server. Jenkins helps to automate the non-human part of the software development process, with continuous integration and facilitating technical aspects of continuous delivery."[2].

It is an easy solution for creating an environment of continuous integration and delivery for many different languages. Jenkins supports different applications for version control such as Git. One of the main reasons for its popularity is its easy installation, configuration and appealing web interface. There is also a large amount of plugins that can be used for Jenkins to work with different tools. Works can also be distributed between different machines, eg. for faster deployment[3].

Jenkins do support both simple jobs and something called pipelines. The latter can configure the whole life-cycle of an application, from commit to delivery, in a very complex way. The pipeline is often described as different blocks that are executed sequentially. These blocks could be named Build, Test and

Deploy and do contain further steps and tasks. Some of the pipeline functionality provided by Jenkins are that pipelines are restored after a system crash, they can be paused, multiple pipelines can work in parallel, checkpoints can be used and more[4][5][6].

2.2 Codeship

Codeship offers different plans and basically handles everything related to CI for the customers[7]. Codeship was acquired by the Cloudbees company in 2018. Before this time there were many different providers of Continuous Integration tools, but in the later years many of them have merged. With a user friendly interface, and provided dependencies it helps to "automate the testing and deployment of the code". It supports fast building by parallelism and caching. Also, the organisation has a strong support organisation. The overall goal is that the customer should focus on their product and code, not on the configuration of continuous integration. [8]

2.3 Travis

Travis is like Codeship, also a hosted CI service. The difference from Jenkins is that the enterprise version isn't free and if more support is needed the price increases. However, the community of Jenkins is larger. It is easier to configure Travis and it has a shorter time to get started. The use of YAML files for configuration makes it possible for advanced configurations. It also has support for Git and parallel execution. So if you need a lot of in depth configuration probably Jenkins is the choice for you. But if you want an easy startup process you should consider Travis. Today Jenkins is the most popular of the two based on questions on Stack Overflow[9].

2.4 Bamboo

Bamboo is a Continuous Integration service created by Atlassian. It was provided both in the cloud and as a local service, but they stopped with the cloud version in 2017. As the other tools it does support automatic build, test and deploy of software as well as parallel jobs. It also does have support for version control like Git and testing frameworks. One of the big differences from Jenkins is that there are more built in support and plugins so there is less need of configuration. It isn't open source as Jenkins and some people could consider it more simple. The pricing is depending on how many build agents you use and this could lead to a better cost for small companies[10].

2.5 On-Premise CI vs Hosted CI

The two different types of CI have their own advantages and disadvantages. With an on-premise CI such as Jenkins, you have complete control of it. The CI can be configured to one's own needs and are not in the mercy of a third-part

service. There is a huge supply of plugins to allow all manner of customisation. There is no need to share resources with other customers of a host based service and from a security point of view, you are in control of it yourself for better or worse. One example is the recent Amazon AWS S3 outage[11]. For all the customization and control one has to divert resources for it. There probably needs to exist someone who knows the system well, staff to be assigned to setup, maintain and configure the on-premise CI. On the contrary, a Hosted CI outsources all the work that comes with on-premise CI to a third party who handles it all. In exchange for control, customization and money. Furthermore, one has to share the resources of the third party and are at mercy of their service. Just as the incident given above, if it were to happen, how are the consequences handled. There are many factors one gives up by choosing a Hosted CI.

3 Microsoft, Amazon and Google and the "Cloud"

With the advent of the cloud, new services and options have opened up to developers. There are many advantages with utilising the cloud and more developers are relying on cloud, a trend that is likely to continue. Even though cloud has been around for over a decade, it isn't until recent years it has seen a massive growth[12]. With the market now seeing the potential with cloud computing.

Many companies use their tools for Continuous Integration on local servers, maintained by their own personnel. This could lead to large effort and costs involved in maintaining the infrastructure. Therefore, many companies have started to use tools for Continuous Integration in the cloud. In this way there can be more focus on the development itself, while the cloud provider handles the infrastructure, security, support and more[13]. There will still exist a need of configuration, but not at the same level and it will mostly be involved with the pipeline itself[14].

"Simply put, cloud computing is the delivery of computing services—including servers, storage, databases, networking, software, analytics, and intelligence—over the Internet ("the cloud") to offer faster innovation, flexible resources, and economies of scale." [15]

If cloud computing frees us from the pain of installing and maintaining physical servers. CI frees us from the pain of building, testing and deploying of the code. Both of the "concepts" aim to remove some of the workload from developers, then it sounds like the two matched together is an excellent choice. Not only does the combination save time, it can also save money.

The global cloud computing market size is expected to grown from \$272 billion in 2018 to \$623 billion by 2023[16]. The Three of the biggest public actors are Amazon Web Services (AWS), Google Cloud and Microsoft Azure with a combined market share of roughly 60%[17].

3.1 AWS

AWS is by far the biggest player right now, and alone has roughly one third of the market share in cloud computing [17]. They offer services such as online storage, virtual cluster of computers to emulate CPU, GPU, RAM etc. databases, web-server, block chain, Internet of Things and much more. Furthermore, they offer DevOps tools such as Continuous Integration in the form of AWS CodePipeline. But also the option to set up other CI such as Jenkins with their services. Currently AWS have support for Cloudbees, Jenkins, Teamcity[18].

AWS CodePipeline is Amazon's continuous delivery service, which is a practice that expands upon CI by developing all code changes to a testing environment after the build stage[19]. AWS CodePipeline runs on their services on the web and has many similarities to a hosted CI, therefore it comes with the advantages associated with hosted CI. However, as of writing it only integrates with one outside source provider, Github[20].

3.2 Google Cloud

Google Cloud Platform offers different services such as Storage, Machine Learning, GPU Computing, Data Analytics, Security, Networking and more. But also services within DevOps and Continuous Integration. Some of the tools that are offered are the Developer Tools which include Cloud Build, Jenkins and Tekton. Cloud build is Google's own tool which offers support for all languages and the designing of custom workflows for build, test and deployment between Virtual machines, Kubernetes, Firebase and more. Some of the differences from Jenkins are that the official support is more extensive and that it has container support for the build environment. Tekton is a open source tool for Continuous Integration and Continuous Delivery, that tries to modernize the field and create an industry standard. In Tekton the different tasks in the pipeline are executed in one or multiple containers, with the help of Kubernetes. This could be done both sequential or parallel and you can increase resources dynamically at runtime. It is even possible to build, test and deploy on different cloud providers. This leads to an increasing flexibility and is one of the reasons for its popularity[21][22][23].

3.3 Microsoft Azure

Azure is a cloud computing service from Microsoft that provides many different services. The product is similar to Google Cloud and Amazon Web Services. Some examples are services for Computing, Storage, Machine Learning, Blockchain, Internet of the things and more. Also the Azure Devops service is provided which was originally released in 2013. It supported the Visual Studio service online together with version control, testing and more. Azure Devops provides Continuous Integration pipelines together with version handling such as Git and support for any language. There is also support for deploying with

Google and Amazon. Parallel execution is supported in Azure as well as a new project called Pandas. Pandas tries to automatically test a pull request at different environments such as Windows, Linux and Mac. The basic pricing is free for open source projects and one free job a month for other projects if you don't want to pay [24][25].

3.4 Comparison

AWS, Google Cloud and Azure offers support for third party tools while at the same time offering their own tools as well. Although there are some limits to using third-party tools such as Jenkins on Google Cloud it should still cover the needs of the majority of users. Overall, the different services are very similar when it comes to handling CI and the decision on what tool to use is more a matter of company culture and taste. CI being such an important part of Devops ensures that the companies will continue to support these services and support more in the future. Since the three cloud provider offer different services and infrastructure outside of the CI spectrum, it is important to look at other services related to ones project. They each have their own strength and weaknesses to take into consideration.

4 Conclusion

Continuous Integration is a well established practice today, with cloud providers offering different solutions. Although there exist some limits, there are many benefits. Given the current trend of cloud computing with more services moving to the cloud. The trends point towards increased options of CI integration with the cloud. It is impossible to say that one option is better than the other. One has to adapt the choices to one's own needs and weigh the pros and cons. A large company with a high need of privacy will maybe think twice before running in the cloud. However, a small company could make use of the benefits if the pricing is advantageous. It is also possible to start out with CI alone and later migrate to the cloud depending on the tool. And future development of cloud might add support for other CI such as Travis CI, but keep in mind that not all CI tools are supported by the three cloud providers covered here. On the other hand, the companies might find it in their interest to not support more CI tools and rather promote their own tools.

References

- [1] P. Duvall, S. Matyas, P. Duvall, and A. Glover, *Continuous Integration: Improving Software Quality and Reducing Risk*. A Martin Fowler signature book, Addison-Wesley, 2007.
- [2] Wikipedia, *Jenkins*. Available at [https://en.wikipedia.org/wiki/Jenkins_\(software\)](https://en.wikipedia.org/wiki/Jenkins_(software)), (2020 – 03 – 20).

- [3] Jenkins, *Jenkins*. Available at <https://jenkins.io/>, (2020-03-30).
- [4] O. Son, Bryant, *Building CI/CD pipelines with Jenkins*. Available at <https://opensource.com/article/19/9/intro-building-cicd-pipelines-jenkins>, (2020-03-30).
- [5] I. Heller, Martin, *What is Jenkins? The CI server explained*. Available at <https://www.infoworld.com/article/3239666/what-is-jenkins-the-ci-server-explained.html>, (2020-03-30).
- [6] Jenkins, *Getting started with pipelines*. Available at <https://jenkins.io/pipeline/getting-started-pipelines/pipeline-defined>, (2020-03-30).
- [7] C. Inc, *Simple CI CD Service Costs: Codeship Basic Pricing*. Available at <https://codeship.com/pricing/basic>, (2020-03-30).
- [8] CloudBees, *Wikipedia CloudBees*. Available at <https://en.wikipedia.org/wiki/CloudBees>, (2020-03-30).
- [9] Travis, *Travis vs Jenkins*. Available at <https://www.guru99.com/jenkins-vs-travis.html>, (2020-03-30).
- [10] Bamboo, *Bamboo*. Available at <https://www.automation-consultants.com/bamboo-vs-jenkins/>, (2020-03-30).
- [11] D. T. Etherington, *Amazon AWS S3 outage is breaking things for a lot of websites and apps*. Available at <https://techcrunch.com/2017/02/28/amazon-aws-s3-outage-is-breaking-things-for-a-lot-of-websites-and-apps/>, (2017-02).
- [12] G2, *Exploring the Future of Cloud Computing in 2020 and Beyond*. Available at <https://learn.g2.com/future-of-cloud-computing>.
- [13] K. Yap and H. B. Singh, *10 things to look for when adopting a cloud CI/CD tool*. Available at <https://bitbucket.org/blog/10-things-to-look-for-when-adopting-a-cloud-ci-cd-tool>, (2019-10).
- [14] M. R. B. Setter, *The Benefits of Continuous Integration in the Cloud - via @codeship*. Available at <https://rollout.io/blog/the-benefits-of-continuous-integration-in-the-cloud/>, (2017-08).

- [15] A. beginner's guide — Microsoft Azure, *What is cloud computing*. Available at <https://azure.microsoft.com/en-in/overview/what-is-cloud-computing/>, (2020-03-30).
- [16] M. R. Firm, *Cloud Computing Market*. Available at <https://www.marketsandmarkets.com/Market-Reports/cloud-computing-market-234.html/>, (2020-03-30).
- [17] F. Richter, Statista Infographics, *Infographic: Amazon Leads 100 Billion Cloud Market*. Available at <https://www.statista.com/chart/18819/worldwide-market-share-of-leading-cloud-infrastructure-service-providers/>, (2020-02).
- [18] AWS, *AWS CodePipeline product integrations*. Available at <https://knapsackpro.com/https://aws.amazon.com/codepipeline/product-integrations/>, (2020-03-30).
- [19] I. A. Freeman, Emily John Wiley Sons, *DevOps*. Available at <https://aws.amazon.com/devops/continuous-delivery/>, (2019).
- [20] S. G. P. Amazon Engdahl, *Blogs*. Available at <https://aws.amazon.com/blogs/devops/integrating-git-with-aws-codepipeline/>, (2008).
- [21] IBM, *Tekton-A-Modern-Approach-to-Continuous-Delivery*. Available at <https://www.ibm.com/cloud/blog/tekton-a-modern-approach-to-continuous-delivery>.
- [22] J. vs Google Cloud Build comparison of Continuous Integration servers, *Jenkins vs Google Cloud Build comparison of Continuous Integration servers What are the differences between Jenkins and Google Cloud Build?* Available at https://knapsackpro.com/ci_comparisons/jenkins/vs/google-cloud-build/, (2020-03-30).
- [23] Google, *Google Cloud Build*. Available at <https://console.cloud.google.com/marketplace/details/google/cloudbuild.googleapis.com>, (2020-03-30).
- [24] Microsoft, *Azure*. Available at <https://azure.microsoft.com/en-us/services/devops/pipelines/>, (2020-03-30).

- [25] Wikipedia, *Wikipedia Azure*. Available at https://en.wikipedia.org/wiki/Microsoft_Azure, (2020 – 03 – 30).